

Introduction to Artificial Intelligence and Applications

Semester One | Course One

Unit 2 : AI Concepts, Terminology and Application Domains

Recap From Unit 1

In the early days of **AI**, the top-down approach to creating intelligent systems. The idea was to extract the knowledge from people into some machine-readable form, and then use it to automatically solve problems. This approach was based on two big ideas:

**Knowledge
Representation**

Reasoning

- **It is important to differentiate knowledge from information or data.**

 **Knowledge** is something which is contained in our head and represents our understanding of the world.

 It is obtained by an active **learning** process, which integrates pieces of information that we receive into our active model of the world.

Recap From Unit 1

- **It is important to differentiate knowledge from information or data.**
- **Data** is something represented in physical media, such as written text or spoken words. Data exists independently of human beings and can be passed between people.
- **Information** is how we interpret data in our head. For example, when we hear the word computer, we have some understanding of what it is.
- **Knowledge** is information being integrated into our world model. For example, once we learn what a computer is, we start having some ideas about how it works, how much it costs, and what it can be used for. This network of interrelated concepts forms our knowledge.
- **Wisdom** is yet one more level of our understanding of the world, and it represents meta-knowledge, eg. some notion on how and when the knowledge should be used.





Image from Wikipedia, By Longlivetheux - Own work, CC BY-SA 4.0

Expert Systems

Expert Systems

- One of the early successes of **AI** were so-called **expert systems**
 - Computer systems that were designed to act as an **expert** in some **limited problem domain**.
- They were based on a **knowledge base** extracted from one or more **human experts**, and they contained an **inference engine** that performed some **reasoning** on top of it.



Expert Systems


Expert systems are built like the human reasoning system, which contains **short-term memory** and **long-term memory**. Similarly, in knowledge-based systems we distinguish the following components:

- **Problem Memory:**
 - Stores the current information about the problem being solved.
 - **Example:** A patient's temperature, blood pressure, or whether they have inflammation.
 - Known as **static knowledge** because it represents what we know right now—a snapshot of the problem.
- **Knowledge Base:**
 - Holds long-term, general knowledge about the problem domain.
 - This knowledge is gathered from experts and remains the same for all consultations.
 - Called **dynamic knowledge** because it helps transition from one state of the problem to another.
- **Inference Engine:**
 - Manages the entire problem-solving process.
 - **Tasks:**
 - Searches through the current problem state.
 - Asks questions to gather missing information.
 - Identifies and applies the correct rules to move forward.

Expert Systems

As an example, let's consider the following expert system of determining an animal based on its physical characteristics:

- This diagram is called an **AND-OR tree**, and it is a graphical representation of a set of production rules
- Drawing a tree is useful at the beginning of extracting knowledge from the expert.
- To represent the knowledge inside the computer it is more convenient to use rules.



[Image by Dmitry Soshnikov](#)

Expert Systems

You can notice that each condition on the of the rule and the action are essentially object-attribute-value (OAV) triplets.

Working memory contains the set of OAV triplets that correspond to the problem currently being solved. A rules engine looks for rules for which a condition is satisfied and applies them, adding another triplet to the working memory.



IF the animal eats meat
OR (animal has sharp teeth
AND animal has claws
AND animal has forward-looking
eyes) THEN the animal is a
carnivore

Expert Systems

Conditions:

- 1.Animal - Eats - Meat
- 2.(If the animal eats meat)
- 3.Animal - Has Teeth - Sharp
- 4.(If the animal has sharp teeth)
- 5.Animal - Has Claws - True
- 6.(If the animal has claws)
- 7.Animal - Has Eyes - Forward-Looking
- 8.(If the animal has forward-looking eyes)

Action:

- 1.Animal - Is - Carnivore
- 2.(Then the animal is a carnivore)



IF the animal eats meat
OR (animal has sharp teeth
AND animal has claws
AND animal has forward-looking
eyes) THEN the animal is a
carnivore

These triplets form the basis for reasoning in the working memory, where conditions are evaluated, and new knowledge (e.g., "Animal - Is - Carnivore") is added when rules are satisfied.

Forward vs. Backward Inference

Forward Inference (Data-Driven Reasoning)

- Starts with what we already know (initial data in working memory).
- Works step by step to add new knowledge until we reach the answer.

Steps:

1. Check if the goal is already known: If yes, stop and return the result.
2. Find rules to apply: Look for rules whose conditions match the current data (this creates a "conflict set").
3. Choose a rule to apply (Conflict Resolution):
 - Pick the first matching rule.
 - Choose randomly.
 - Pick the rule with the most specific conditions.
4. Apply the rule: Use the rule to add new knowledge to the working memory.
5. Repeat: Go back to step 1 until the goal is reached.

Example: Diagnosing a patient based on existing lab results.

Backward Inference (Goal-Driven Reasoning)

- Starts with a specific goal (e.g., "What is the diagnosis?").
- Asks targeted questions to find the information needed to reach the goal.

Steps:

1. Identify rules for the goal: Find rules that have the desired answer on the Right-Hand Side (RHS).
2. If no rules exist, ask for missing data: If no rules cover the goal, ask the user for the information.
3. Test one rule as a hypothesis: Try to prove the rule by checking its conditions (Left-Hand Side, LHS).
4. Repeat for sub-goals: If the rule's conditions depend on other unknowns, repeat the process for those sub-goals.
5. If a rule fails: Try another rule from step 3.

Example: In medical diagnosis, asking for specific tests or symptoms only when needed to narrow down the cause.

Implementing Expert Systems

Expert systems can be implemented using different tools:

- Programming them directly in some **high level programming language**. This is not the best idea, because the main advantage of a knowledge-based system is that knowledge is separated from inference, and potentially a problem domain expert should be able to write rules without understanding the details of the inference process
- Using **expert systems** shell, i.e. a system specifically designed to be populated by knowledge using some knowledge representation language.

Let us implement some code



Ontologies and the Semantic Web

Background:

- By the end of the 20th century, there was a push to improve how we find information on the Internet.

The idea:

- Use knowledge representation to annotate web resources so users can make very specific queries.
- This idea became known as the **Semantic Web**.

Smart Knowledge Representation

- A way to describe knowledge using **Description Logics (DL)**, which provide a formal, structured, and logical framework for understanding data.
- Helps in organizing information into hierarchies and assigning properties to objects (like categories and subcategories).
- Beyond organizing, DL allows systems to reason or infer new facts automatically.
 - **Ontology about transportation:**
 - "All cars are vehicles."
 - "Toyota Camry is a car."
 - The system infers: "Toyota Camry is a vehicle."

Ontologies and the Semantic Web

Is it usable today?

- Yes. Modern AI and search engines use similar techniques to infer relationships and provide better recommendations or search results.
 - **Google** uses ontologies to link related search terms (e.g., searching "Canine" may show results for "Dog").

Distributed Knowledge

- Knowledge is distributed across multiple systems or websites, but concepts are uniquely identified using global **URIs (Uniform Resource Identifiers)**.
A "Dog" ontology in one database and a "Veterinary Care" ontology in another **can link information seamlessly**.
- Makes data interoperable (**easily shared and understood across systems**). _____.
- **Enables linking knowledge** from diverse domains—biology, transportation, healthcare, etc.—on the global web.
- _____
- _____

Ontologies and the Semantic Web

Special XML-Based Languages

- **Tools to describe, structure, and share knowledge in a machine-readable format:**
 - RDF (Resource Description Framework): A standard way to represent information about resources.
 - Example: Describe "Paris" with RDF:
 - "Paris is a city. It is in France."
 - RDFS (RDF Schema): Adds hierarchies and relationships.
 - Example: Define "City" as a subclass of "Location."
 - OWL (Ontology Web Language): More expressive, allowing for rules and reasoning.
 - Example: Define "City" and infer, "If Paris is a city, it is also a location."
- **Is it usable today?**
 - Yes. These languages are actively used in:
 - **Knowledge Graphs:** Google, Microsoft, and Amazon build these to link data intelligently.
 - **AI and Research:** Fields like healthcare use OWL to create ontologies for diseases and treatments.

Core Concept: Ontology

What is an Ontology?

- A formal specification of a domain (a specific topic or area).
- **Organizes knowledge as:**
 - A hierarchy of objects (simple version).
 - Rules and relationships for inference (advanced version).
- **Purpose of Ontologies in the Semantic Web:**
 - Help machines understand the meaning of data.
 - Example: In an ontology about animals:
 - Define "Dog" as a type of "Mammal."
 - Add rules: "Mammals have hair" → Automatically know "Dogs have hair."

Why is This Important?

- Enables the web to move beyond simple keyword searches to understanding user queries.
- Makes the web smarter, linking knowledge across sites.
- Supports applications like intelligent search engines, AI assistants, and data integration tools.

Core Concept: Ontology

Real-World Applications of Ontologies in Food and Nutrition:

- **Health Apps:**
 - An ontology can link food items to their nutritional content.
 - Query: "Which foods are rich in Vitamin C?"
 - The system uses the ontology to find and suggest items like oranges, apples, and strawberries.
- **E-commerce:**
 - When you search for "Fruits" on a grocery website, it can automatically group apples, bananas, and other fruits without needing manual categorization.

Let's Code !!

Practise Quiz !!

~~<https://forms.gle/wSSmQejDyq5G9dae7>~~

